Shifting Left on Production Excellence with Observability

DevOps Enterprise Summit
October 5-7, 2021

featuring illustrations
by @emilywithcurls!

Your business wants to move faster

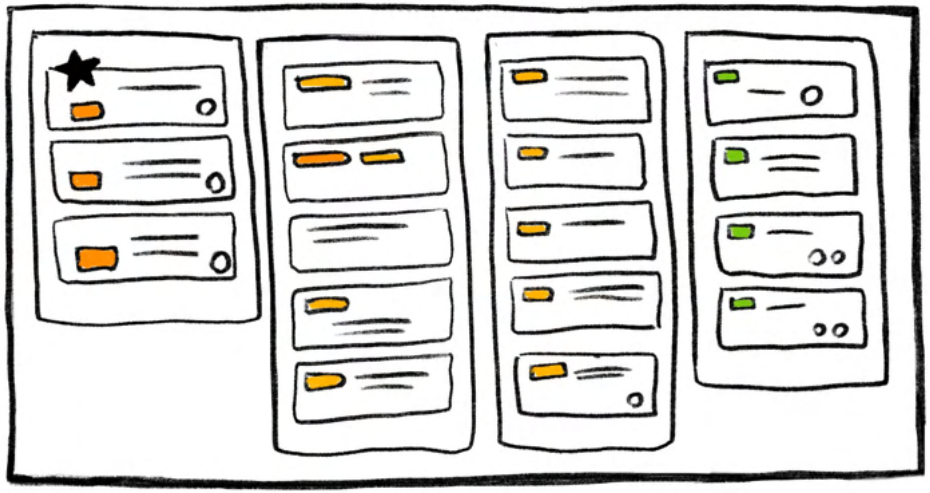

Velocity requires tight feedback loops

We need to shift production left.

Liz Fong-Jones

Principal Developer Advocate at Honeycomb.io

lizthegrey.com

@lizthegrey

Shelby Spees

Site Reliability Engineer at Equinix Metal

spees.dev

@shelbyspees

We've come a long way

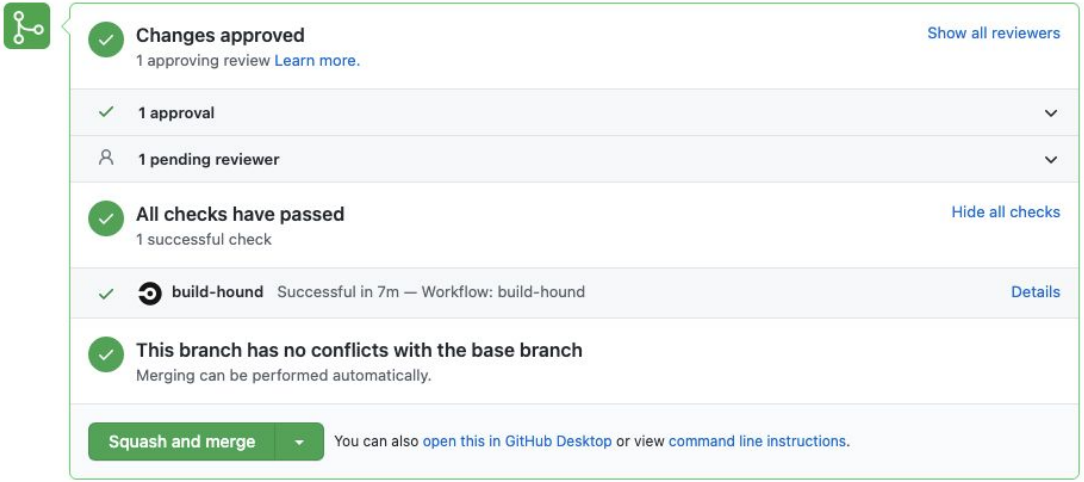


```
# and private subnet for that list index,
# then create an EIP and attach a nat_gateway for each one.  and an aws route
# table should be created for each private subnet, and add the correct nat_gw

resource "aws_subnet" "private" {
  vpc_id            = aws_vpc.mod.id
  cidr_block        = var.private_ranges[count.index]
  availability_zone = var.azs[count.index]
  count             = length(var.private_ranges)
  tags = {
    Name = "${var.env}_private_${count.index}"
  }
}

resource "aws_subnet" "public" {
  vpc_id            = aws_vpc.mod.id
  cidr_block        = var.public_ranges[count.index]
  availability_zone = var.azs[count.index]
  count             = length(var.public_ranges)
  tags = {
    Name = "${var.env}_public_${count.index}"
  }
  map_public_ip_on_launch = true
}

# refactor to take all the route {} sections out of routing tables,
# and turn them into associated aws_route resources
# so we can add vpc peering routes from specific environments.
resource "aws_route_table" "public" {
  vpc_id = aws_vpc.mod.id
  tags = {
    Name = "${var.env}_public_subnet_route_table"
```

What's the next step?

Should developers be on-call?

Production is increasingly complex

Production feels intimidating

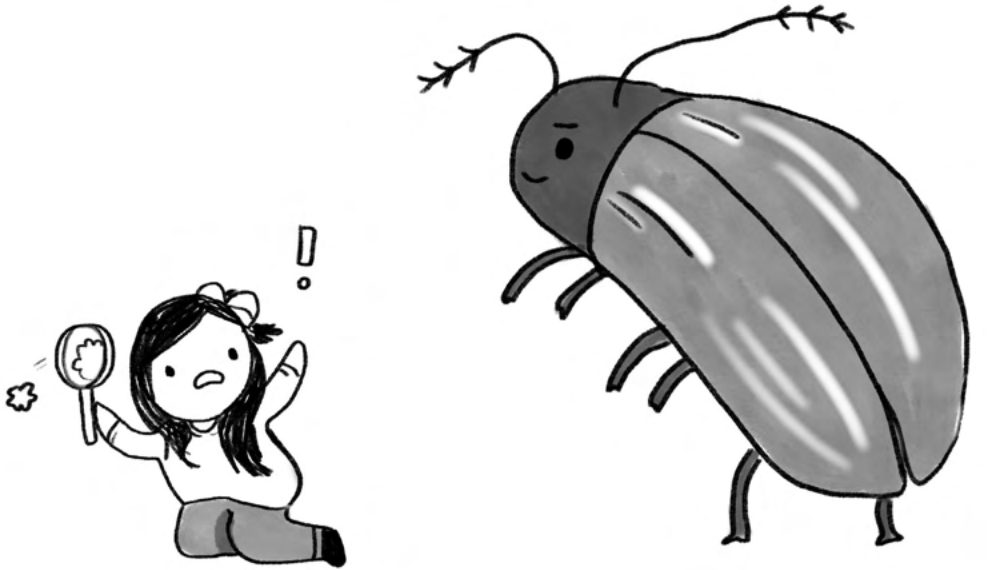

Traditional tools are inscrutable

Prod is always encountering new issues

emergent failure modes

small, unrelated failures cascading together to degrade or take down a system

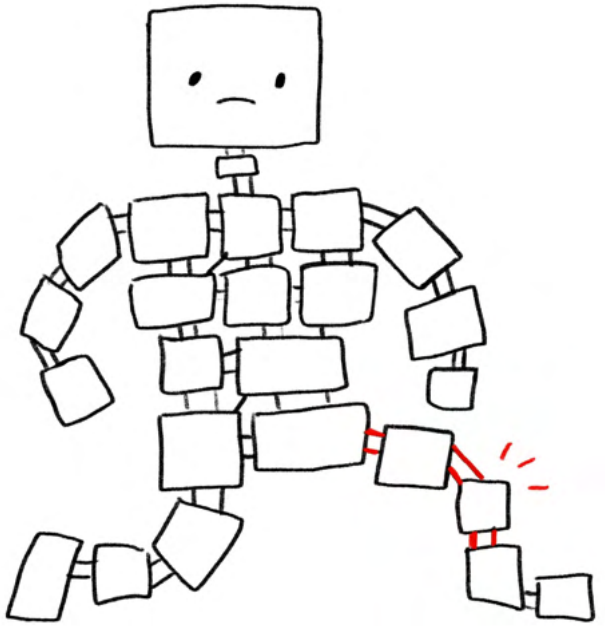

see also: [how.complexsystems.fail]

42%

of developer time is spent dealing with bad code and tech debt

Source: The Developer Coefficient, Stripe, 2018

Teams can’t make forward progress

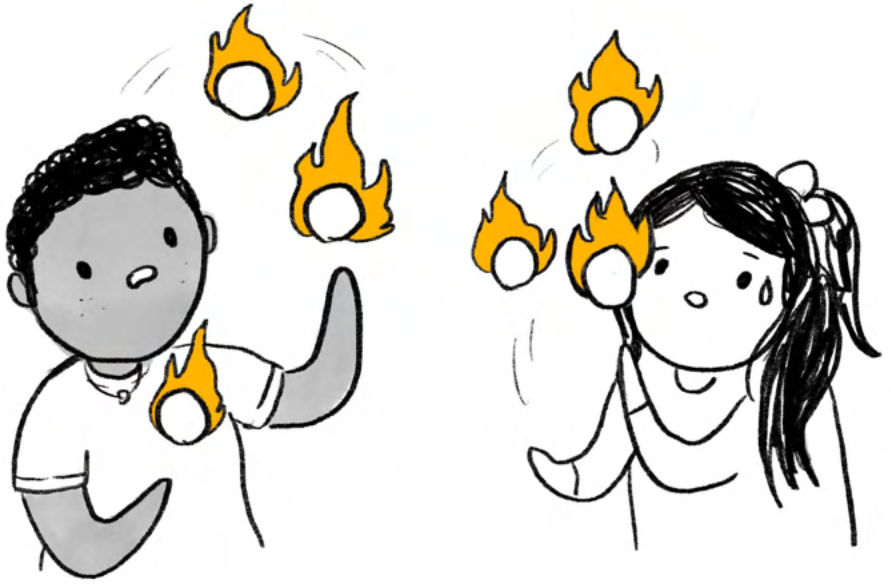

Our heroes are exhausted

We shouldn’t stop here

All teams need production excellence

Can't we just buy it?

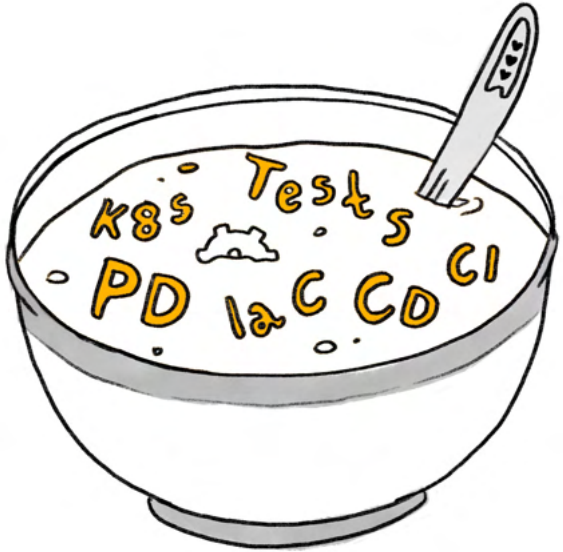

Invest in people, culture, and process

Production Excellence is Business Excellence

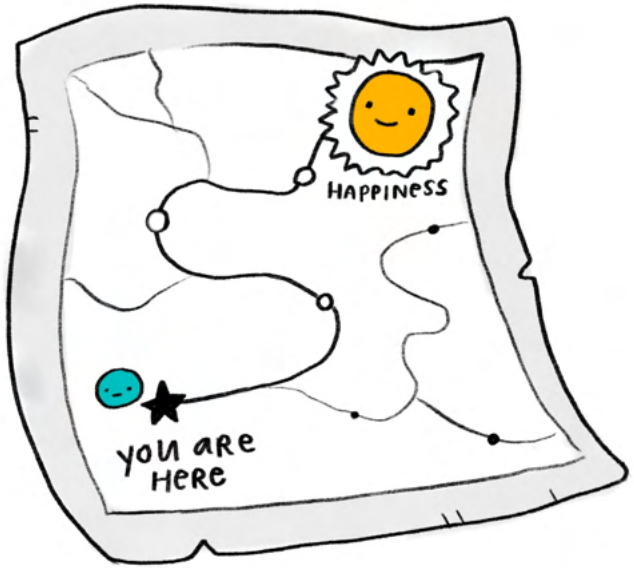

Start with observability

What is observability?

The ability to inspect and understand a system's internal state using the telemetry data it's already outputting.

(Even if your system is in a state you didn't know was possible!)

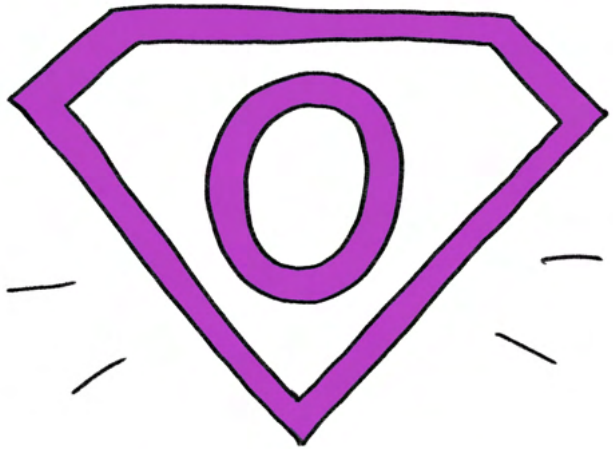

Hard-to-debug problems

Distributed Systems
small change causing downstream effects?

Poor Performance
what is worth optimizing?

Subset of Traffic
only some users are complaining?

Regulatory Requirements
can't touch the box because of SOX?

Observability finds answers.

Distributed Systems
Not a problem, we have traces.

Poor Performance
Identify the bottleneck.

Subset of Traffic
Slice and dice with ad-hoc queries.

Regulatory Requirements
Read-only access.

Example event

```
event = {
    "duration_ms" = 1198,
    "request.id" = 12345,
    "request.method" = "GET",
    "request.path" = "/search",
    "request.query_string" = "category=decor&price<=50",
    "response.status_code" = 200
}
```

Example event

```
event = {
    "trace_id" = 0xC001,
    "span_id" = 12345,
    "parent_id" = 12344,
    "duration_ms" = 1198,
    "request.method" = "GET",
    "request.path" = "/search",
    "request.query_string" = "category=decor&price<=50",
    "response.status_code" = 200
}
```

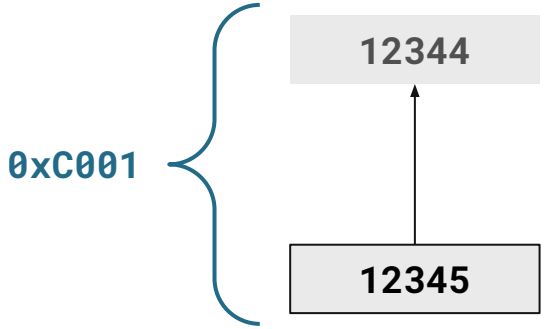

Examples of dimensions to capture

infrastructure:

- build ID
- kubernetes pod
- kafka broker

application:

- library version
- user-agent
- sanitized SQL query string

domain:

- number of items in cart
- coupon code
- payment processor
  (e.g. PayPal vs. Stripe)

Technical Decisions are Business Decisions

How do we map effort to impact?

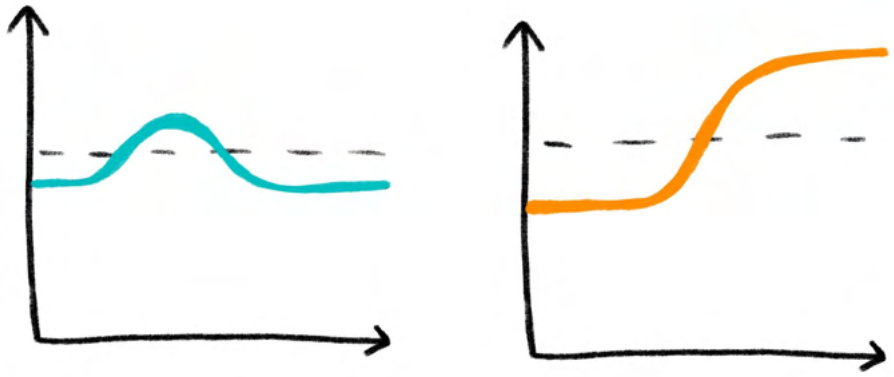

“How does this drive the business forward?”

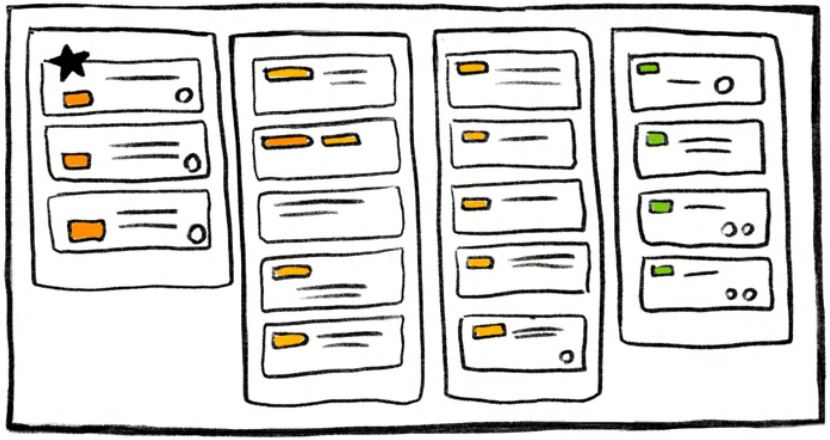

Service Level Objectives (SLOs)

Common language for engineers and business stakeholders

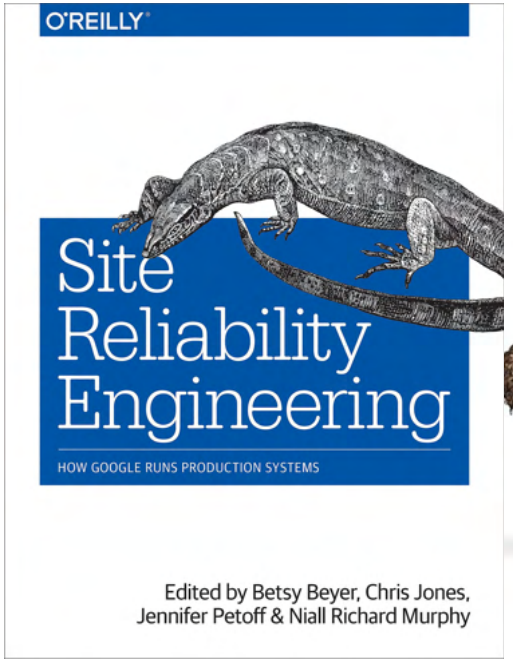

Think in terms of events in context

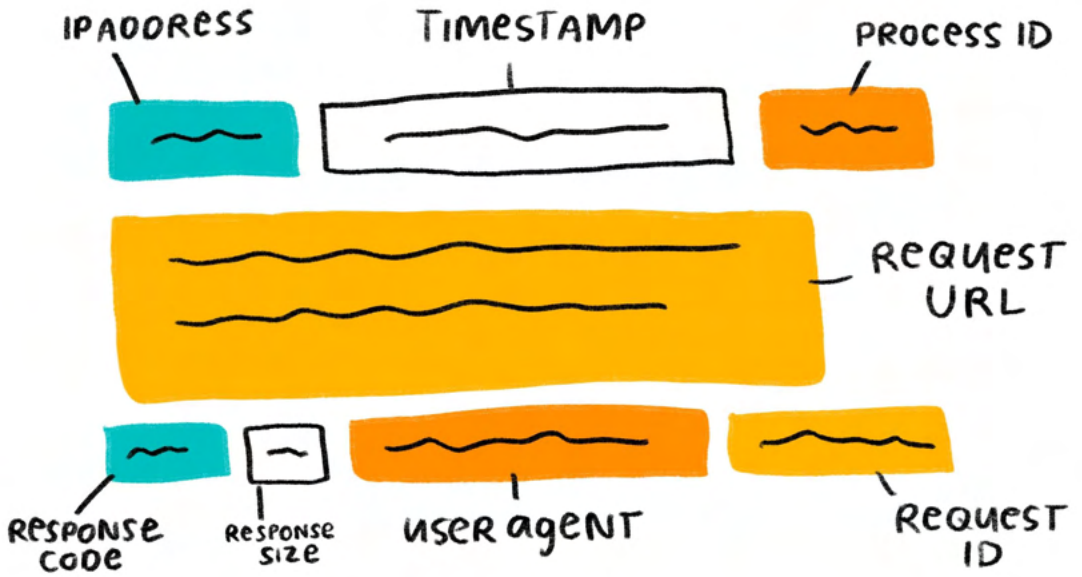

SLI: for each event, good or bad?

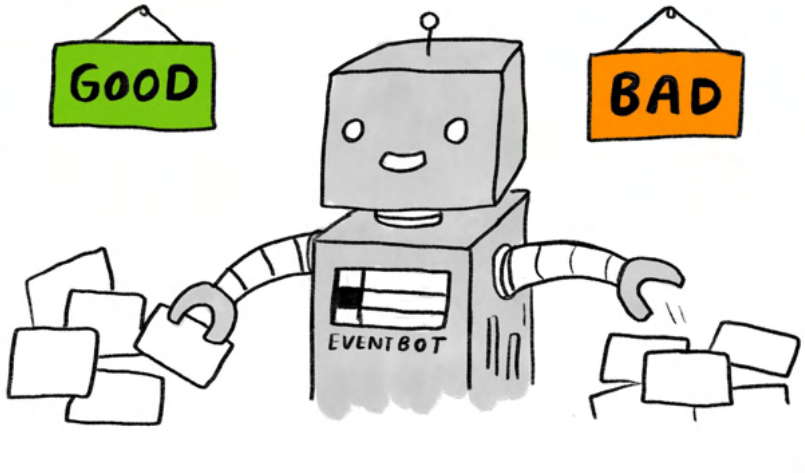

Use a window and target percentage

window is usually 30, 60, or 90 days

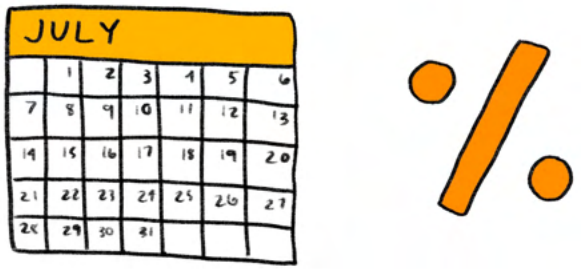

99.9%

of app Home Page loads
over the past 30 days
were “successful” and “fast enough”

Historical SLO compliance

A good SLO barely keeps users happy

Error budget: allowed unavailability

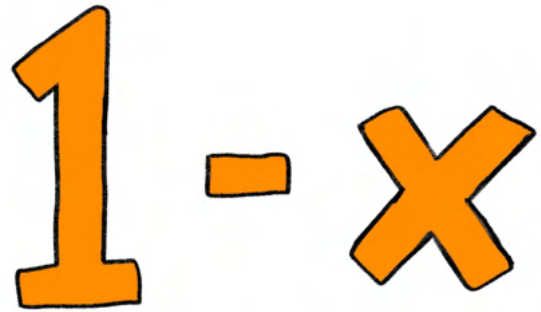

99.9% SLO Target → 0.1% Error Budget

Example error budget

99.9% SLO Target → 0.1% Error Budget

```
      1,000,000   requests/month
×           0.1%
-------------------------------
          1,000   requests/month
```

We're allowed 1000 "bad" requests/month

Error budget remaining

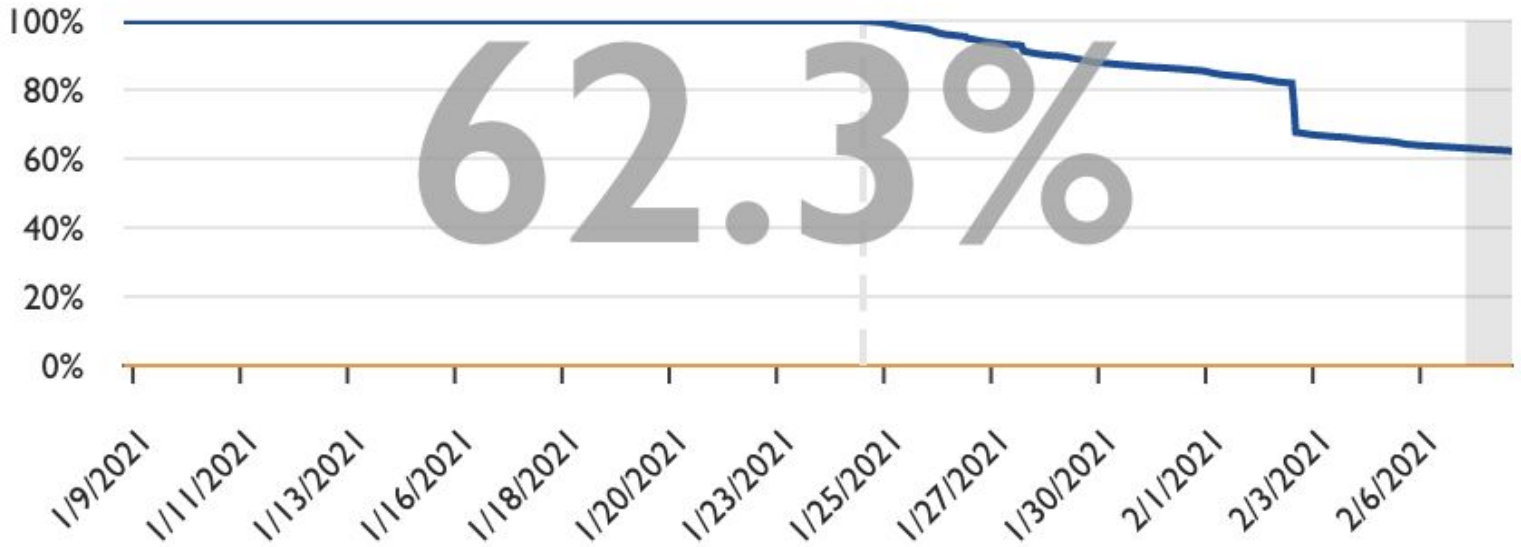

When is it okay to take risks?

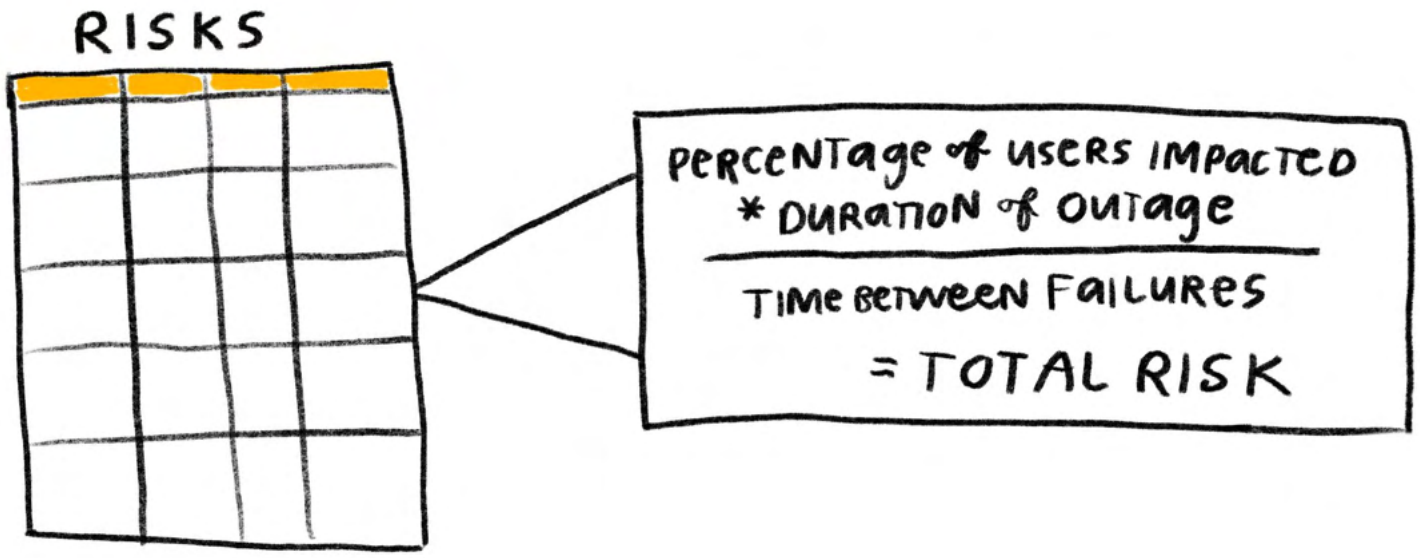

When is it *not* okay?

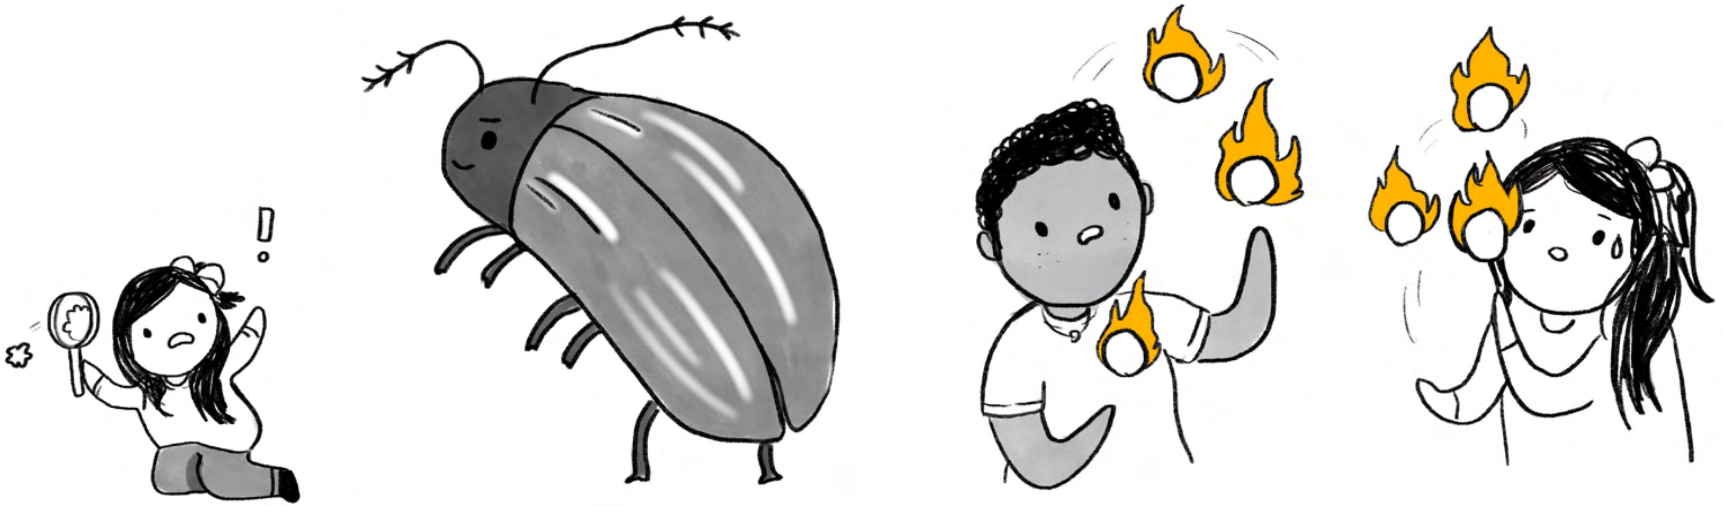

Alerting on SLOs

Alert on error budget **burndown rate**

> “How soon are we going to run out of budget?”

NOT: every **potential cause** of an alert

> “Disk is at 90.05%!!!” (does it matter?)

Observability & Continuous Delivery

@lizthegrey + @shelbyspees

Continuous Delivery (at Honeycomb)

@lizthegrey + @shelbyspees

Instrument as we code.

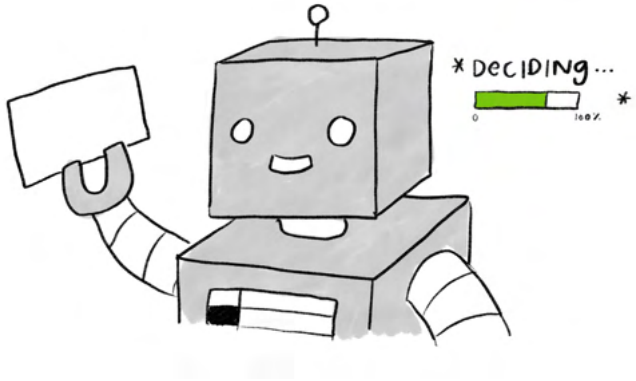

Ensure fast integration.

Observe behavior in prod.

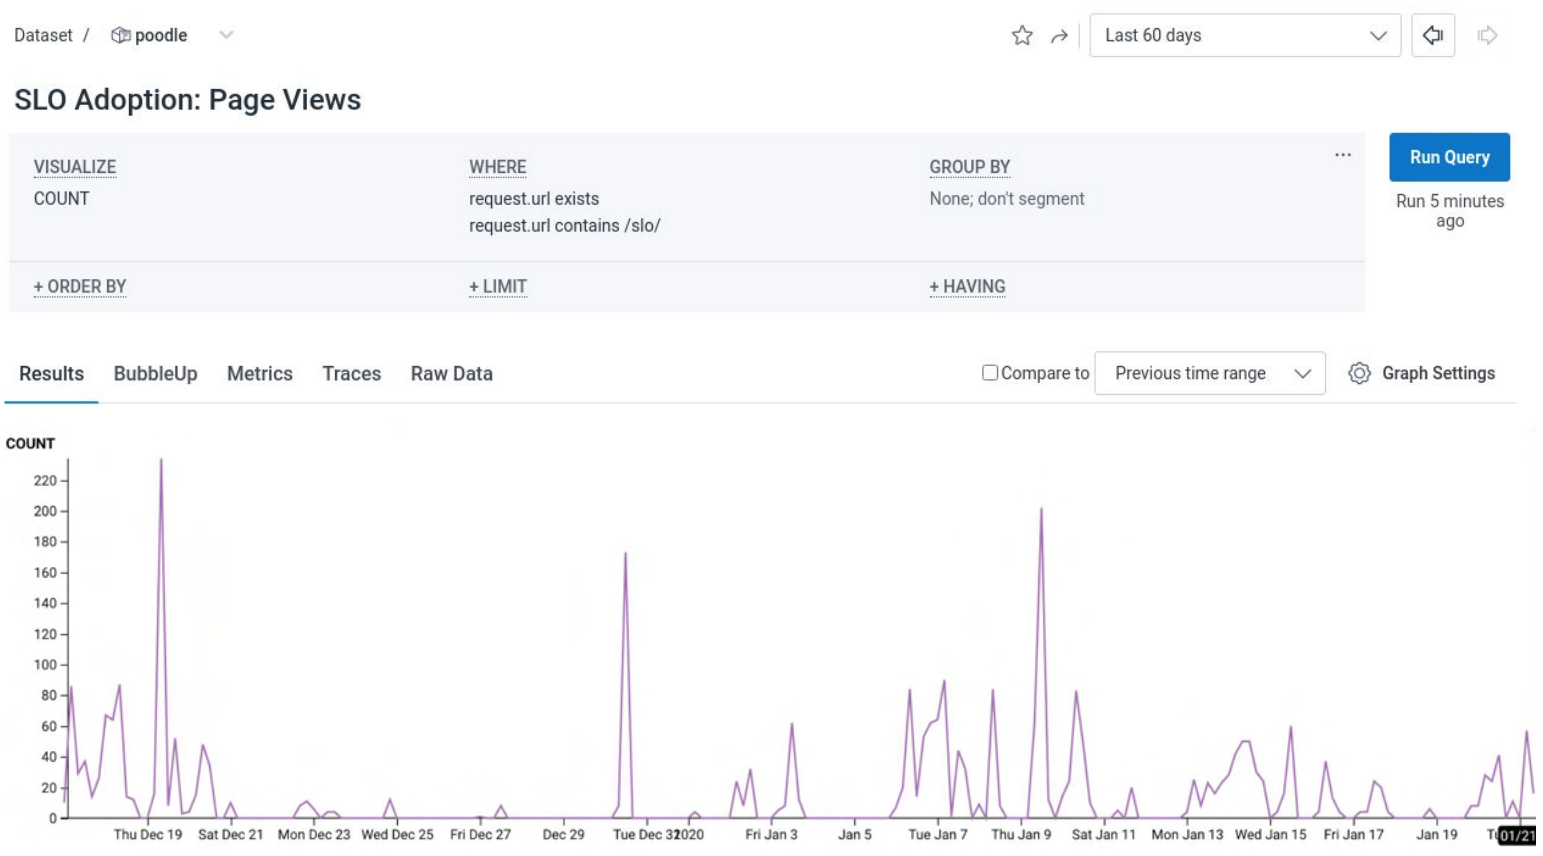
Dataset / poodle
Last 60 days
SLO Adoption: Page Views
VISUALIZE
COUNT
WHERE
request.url exists
request.url contains /slo/
GROUP BY
None; don't segment
Run Query
Run 5 minutes ago
+ ORDER BY
+ LIMIT
+ HAVING
Results
BubbleUp
Metrics
Traces
Raw Data
Compare to
Previous time range
Graph Settings
COUNT
220
200
180
160
140
120
100
80
60
40
20
0
Thu Dec 19
Sat Dec 21
Mon Dec 23
Wed Dec 25
Fri Dec 27
Dec 29
Tue Dec 31
2020
Fri Jan 3
Jan 5
Tue Jan 7
Thu Jan 9
Sat Jan 11
Mon Jan 13
Wed Jan 15
Fri Jan 17
Jan 19
01/21

Bring production excellence to your org

Invest in the right places

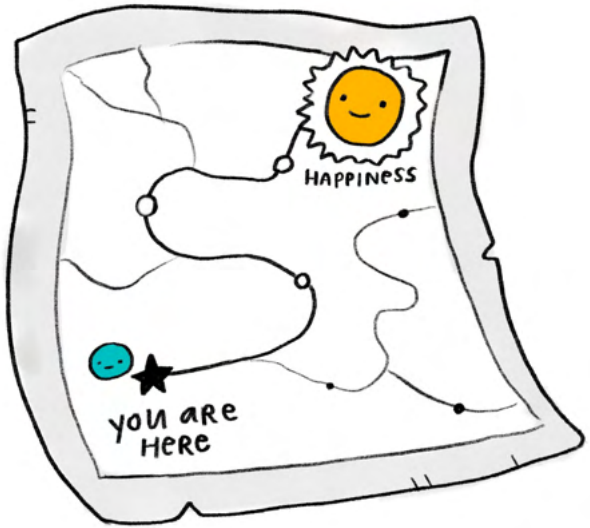

Lean on auto-instrumentation

OpenTelemetry

Instrument your builds

Find a champion

Learn and iterate

SLOs aren't one-and-done

- Adjusting SLO targets
- Improving instrumentation
- Updating alerting methods

see also: go.hny.co/biz-goals-slos

Make it scale to your org

Prioritize:

- developer experience
- on-call onboarding

Invest in:

- config as code
- custom instrumentation libraries

Case Study: Observability at Vanguard

Observability champion(s) ✔️

Prioritizing knowledge transfer ✔️

OpenTelemetry ✔️

Service Level Objectives ✔️

18 million developers

It’s on us to grow them into production engineers.

Help we're looking for

CNCF Observability
Technical Advisory Group (TAG)

- CNCF Slack **#tag-observability** channel

OpenTelemetry

- CNCF Slack **#opentelemetry** channel
- OpenTelemetry.io

OpenSLO

- OpenSLO Slack
- OpenSLO.com

Schedule 30 minutes

- hny.co/meet/liz

Find us on Twitter

- twitter.com/lizthegrey
- twitter.com/shelbyspees